



**POLITECNICO**  
MILANO 1863

# Hardware Architectures for Embedded and Edge AI

Prof Manuel Roveri – [manuel.roveri@polimi.it](mailto:manuel.roveri@polimi.it)

Massimo Pavan – [massimo.pavan@polimi.it](mailto:massimo.pavan@polimi.it)

*Exercise session 4 – TinyML and TensorFlow lite (for Microcontrollers)*

# What is TinyML?

- Fast growing field of Machine Learning
- Algorithms, hardware, and software
- **On-device** sensor data analytics
- Extreme **low power** consumption
- **Always on** ML use-cases
- **Battery operated** devices

# What are the goals of tinyML?

- We want to **perform inference** on an embedded/iot device
- We want to be able to perform computation completely **on-device**, for efficiency, privacy and latency reasons.
- We want it to solve **simple tasks**, with respect to big ML pipelines/algorithms
- We want TinyML algorithms to have a **low power consumption**, in order to be able to function continuously for days, weeks or even a year without human aid and without changing batteries
- Very often, we want it to be «embeddable» in larger cascade architectures.

# What are the application of tinyML now...



Wake-word detection



Presence detection



Anomaly detection

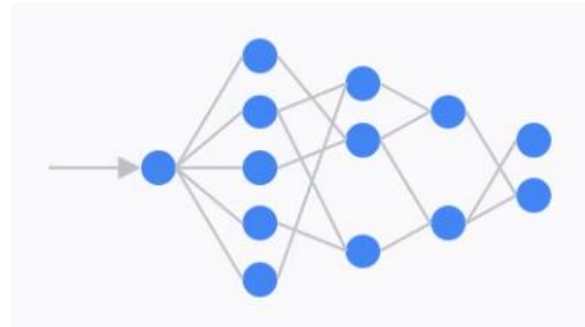
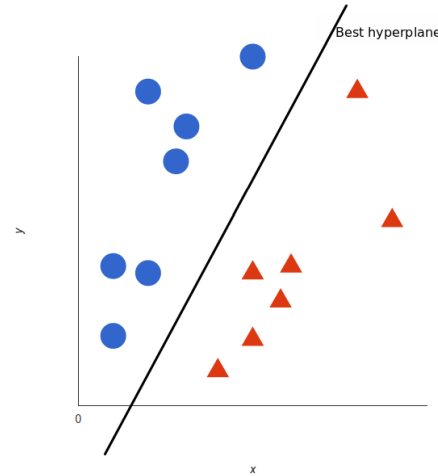
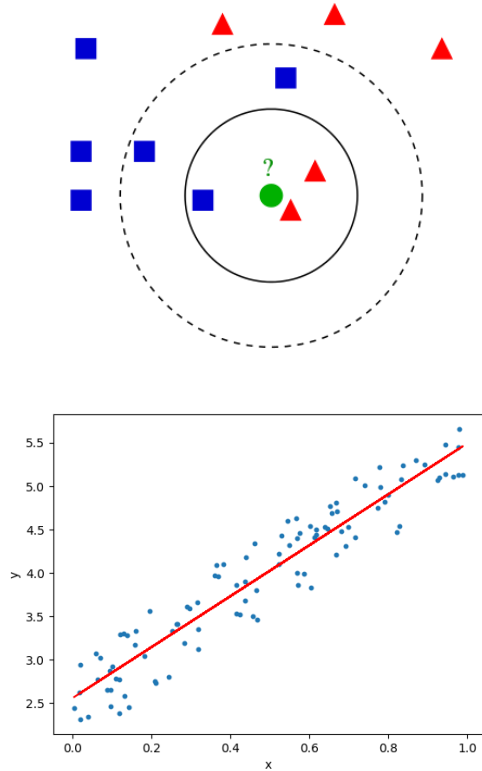


Health parameter monitoring

... and what will they be in the next future



# TinyML is not only Neural Networks, nor only TFmicro



 **TensorFlow** Lite (Micro)

 **CMSIS** **CMSIS-NN**  
CMSIS NN Software Library

   
**STM32** **Cube.AI**

# Endpoints have sensors, a lot of sensors

## **Motion sensors**

Gyroscope, Radar,  
Accelerometer

## **Acoustic sensors**

Ultrasonic, Microphones,  
Vibrometers ...

## **Environmental sensors**

Temperature, Humidity,  
Pressure, IR ...

## **Touchscreen sensors**

Capacitive, IR

## **Image sensors**

Thermal, Image

## **Biometric sensors**

Fingerprint, Heart rate ...

## **Force sensors**

Pressure, Strain

## **Rotation sensors**

Encoders

# Endpoints have sensors, a lot of sensors

## **Motion sensors**

Gyroscope, Radar,  
**Accelerometer**

## **Acoustic sensors**

Ultrasonic, **Microphones**,  
Vibrometers ...

## **Environmental sensors**

Temperature, Humidity,  
Pressure, IR ...

## **Touchscreen sensors**

Capacitive, IR

## **Image sensors**

Thermal, **Image**

## **Biometric sensors**

Fingerprint, Heart rate ...

## **Force sensors**

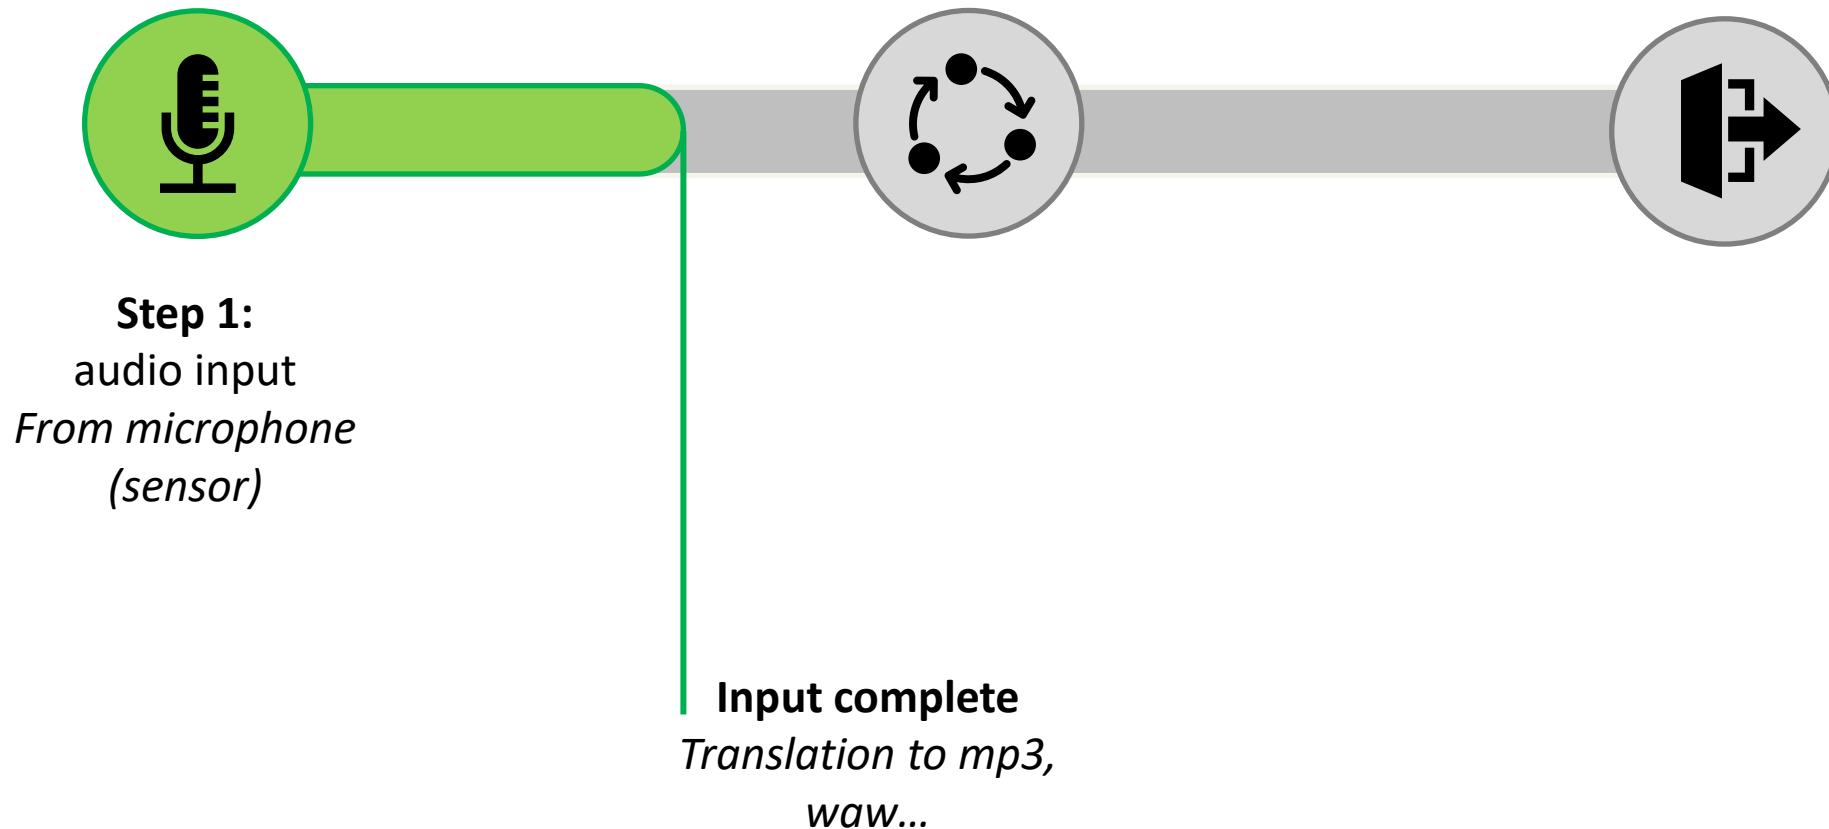
Pressure, Strain

## **Rotation sensors**

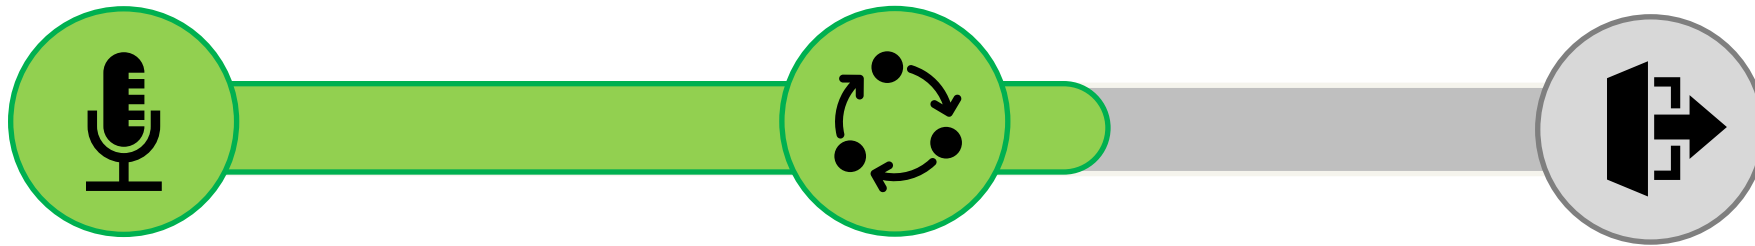
Encoders



# A complete ML pipeline: wake-word detection example

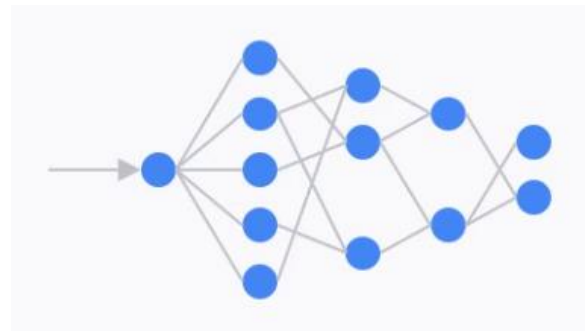


# A complete ML pipeline: wake-word detection example

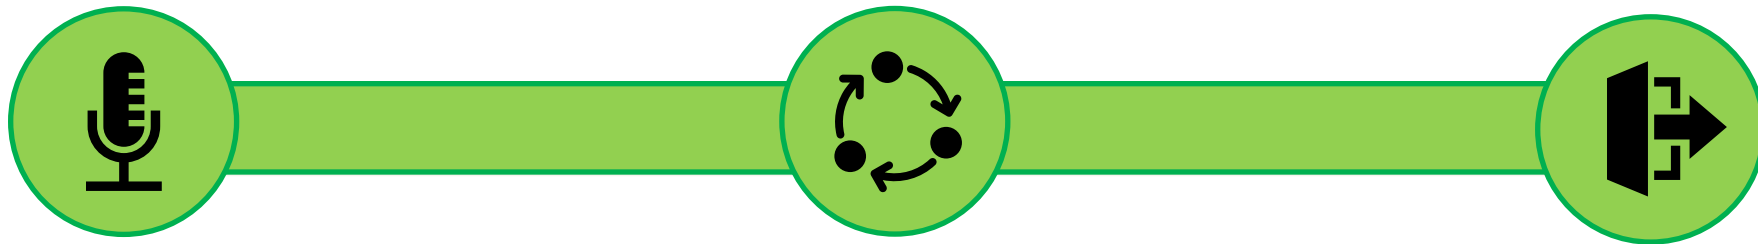


## Step 2:

Process input  
*Translation, then  
execute command*



# A complete ML pipeline: wake-word detection example

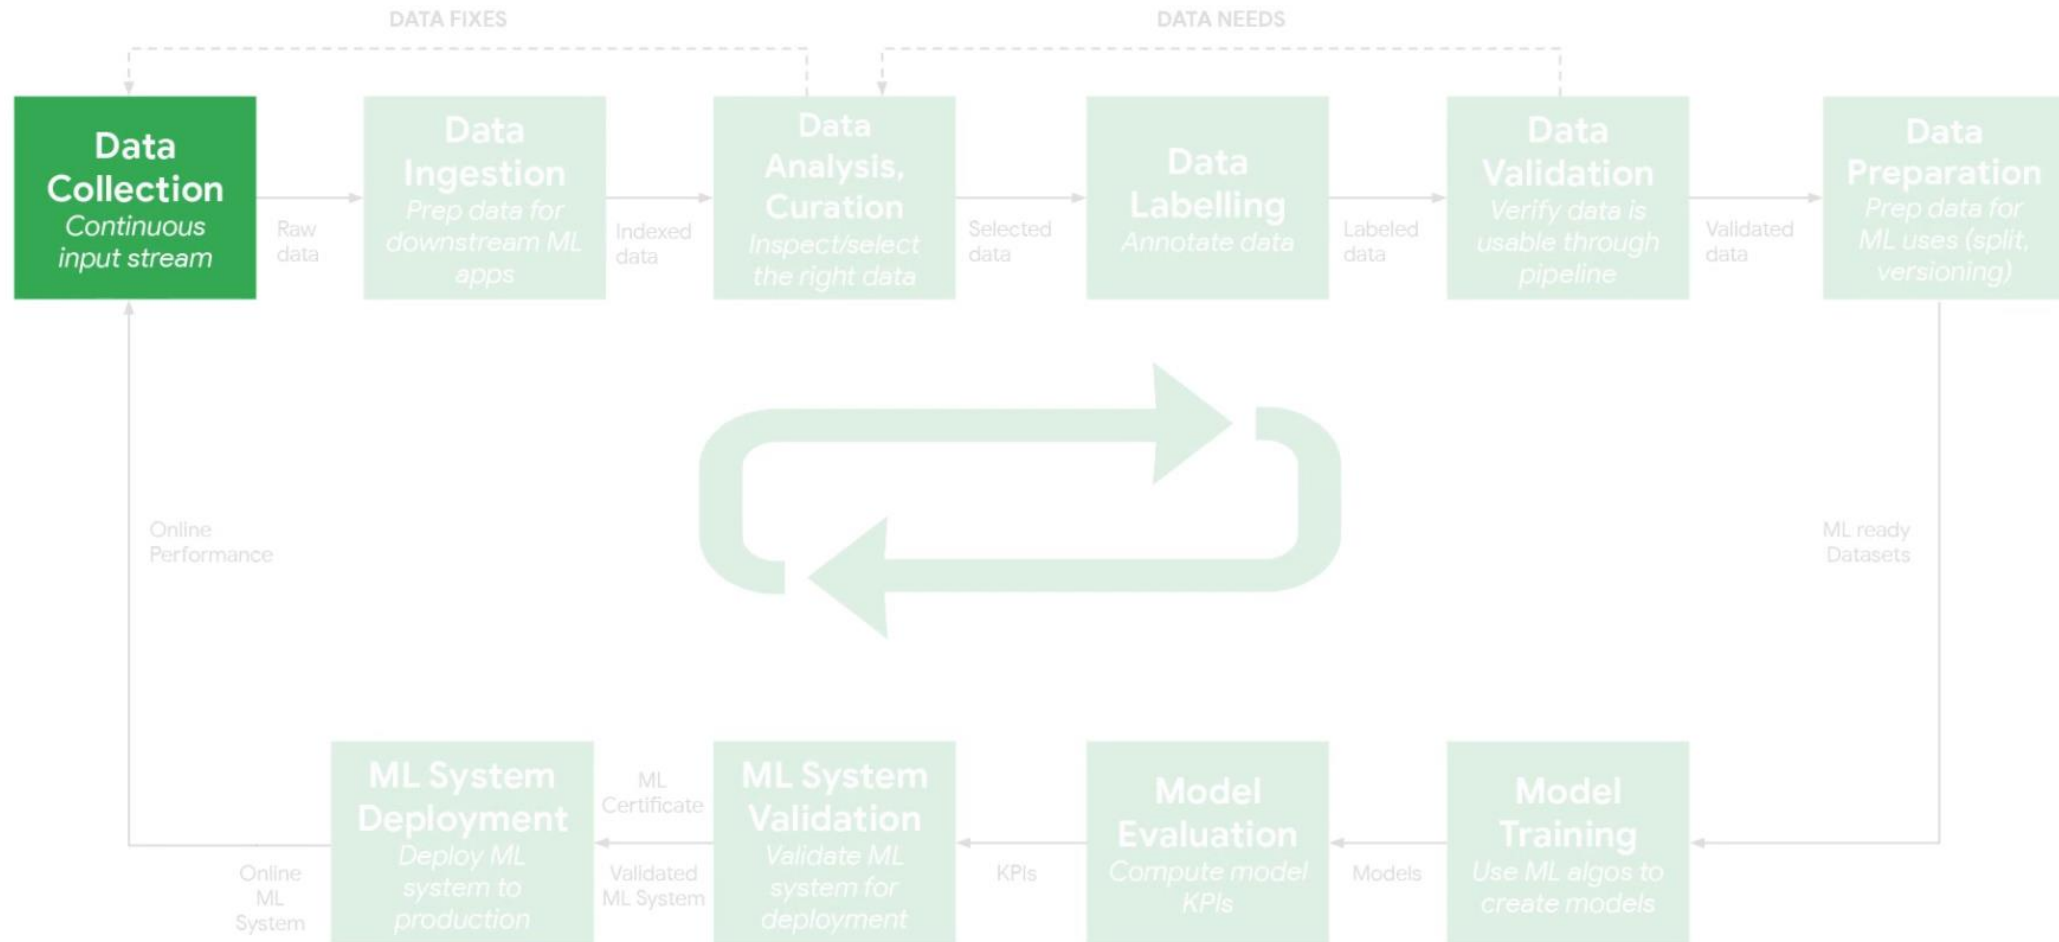


**Step 3:**  
Generate output  
*Play response trough  
embedded speakers*

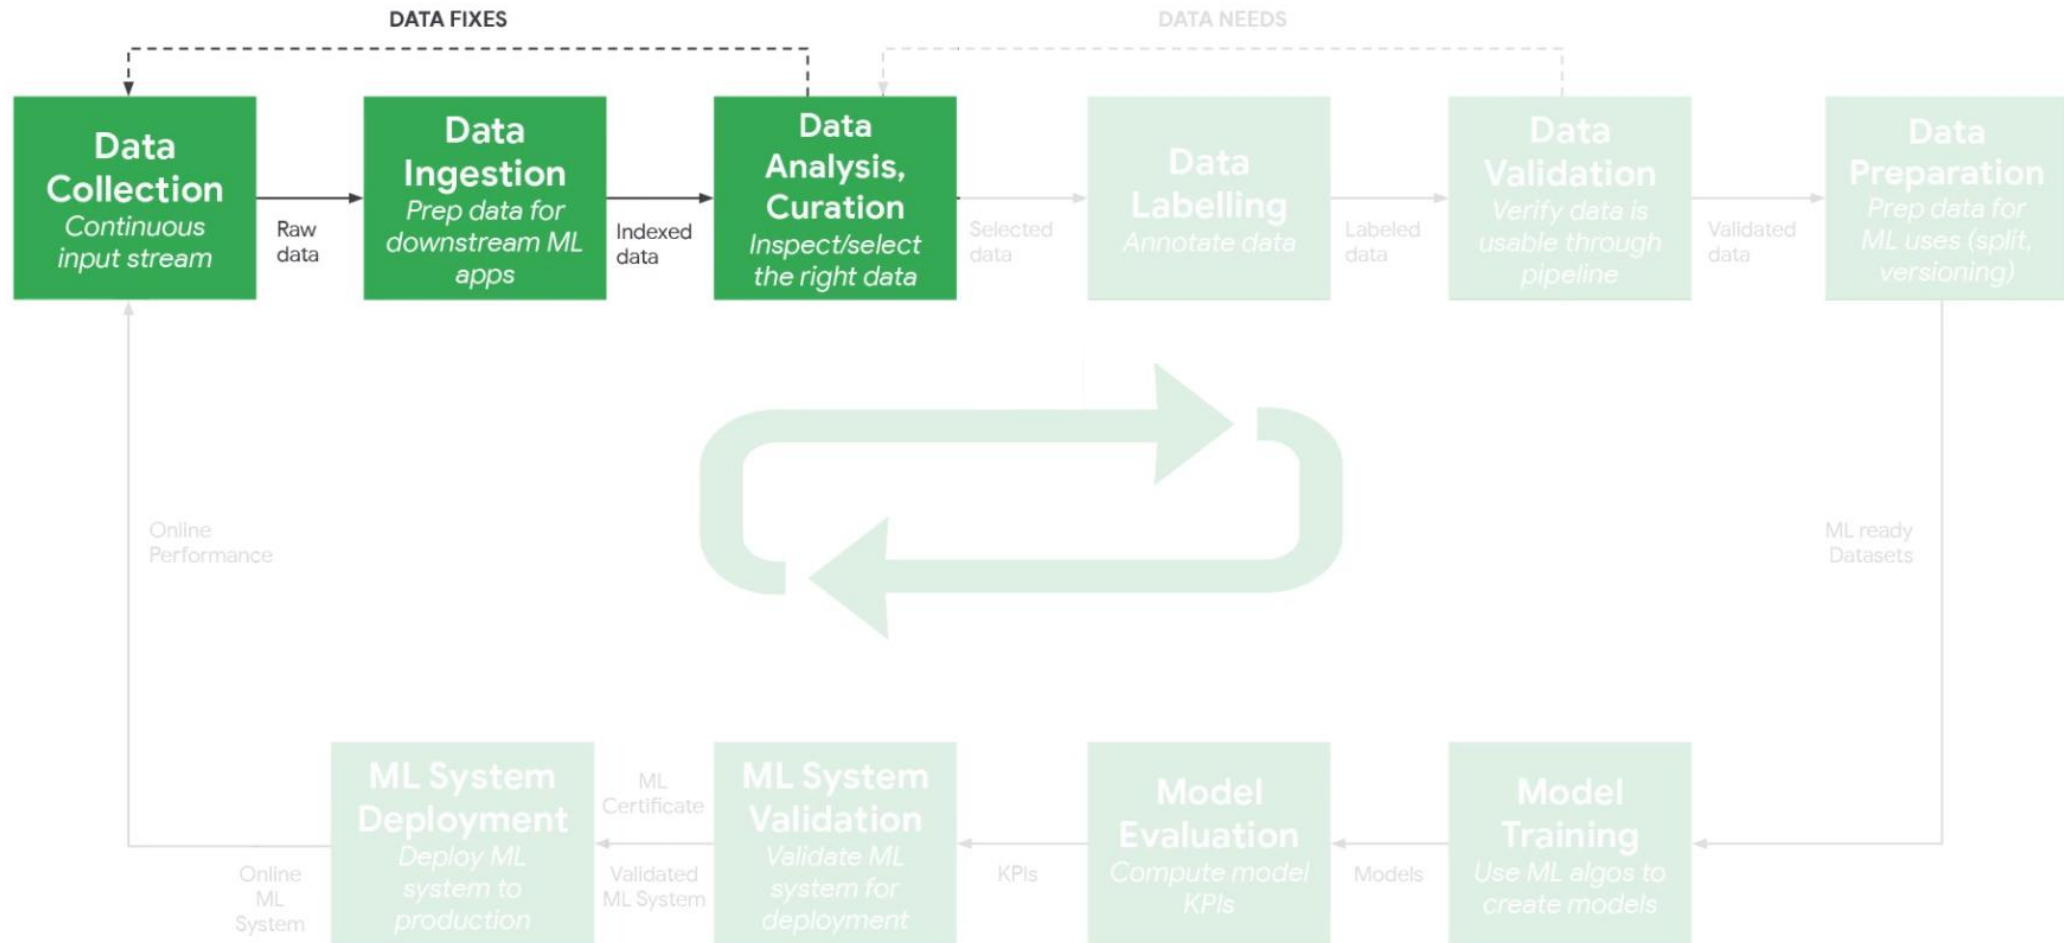
# Questions

- How do we **capture** the data to feed into the network?
- How do you **design** the neural network to take in the speech signal?
- What **dataset** does the network need to be trained on?
- How do we **pre-process** the data for neural network inference?
- How do you **post-process** the neural network output?
- How do you **deploy** this on the microcontroller?
- How do we ensure that the neural network is **resilient**?
- How do we get the neural network to **train faster**?

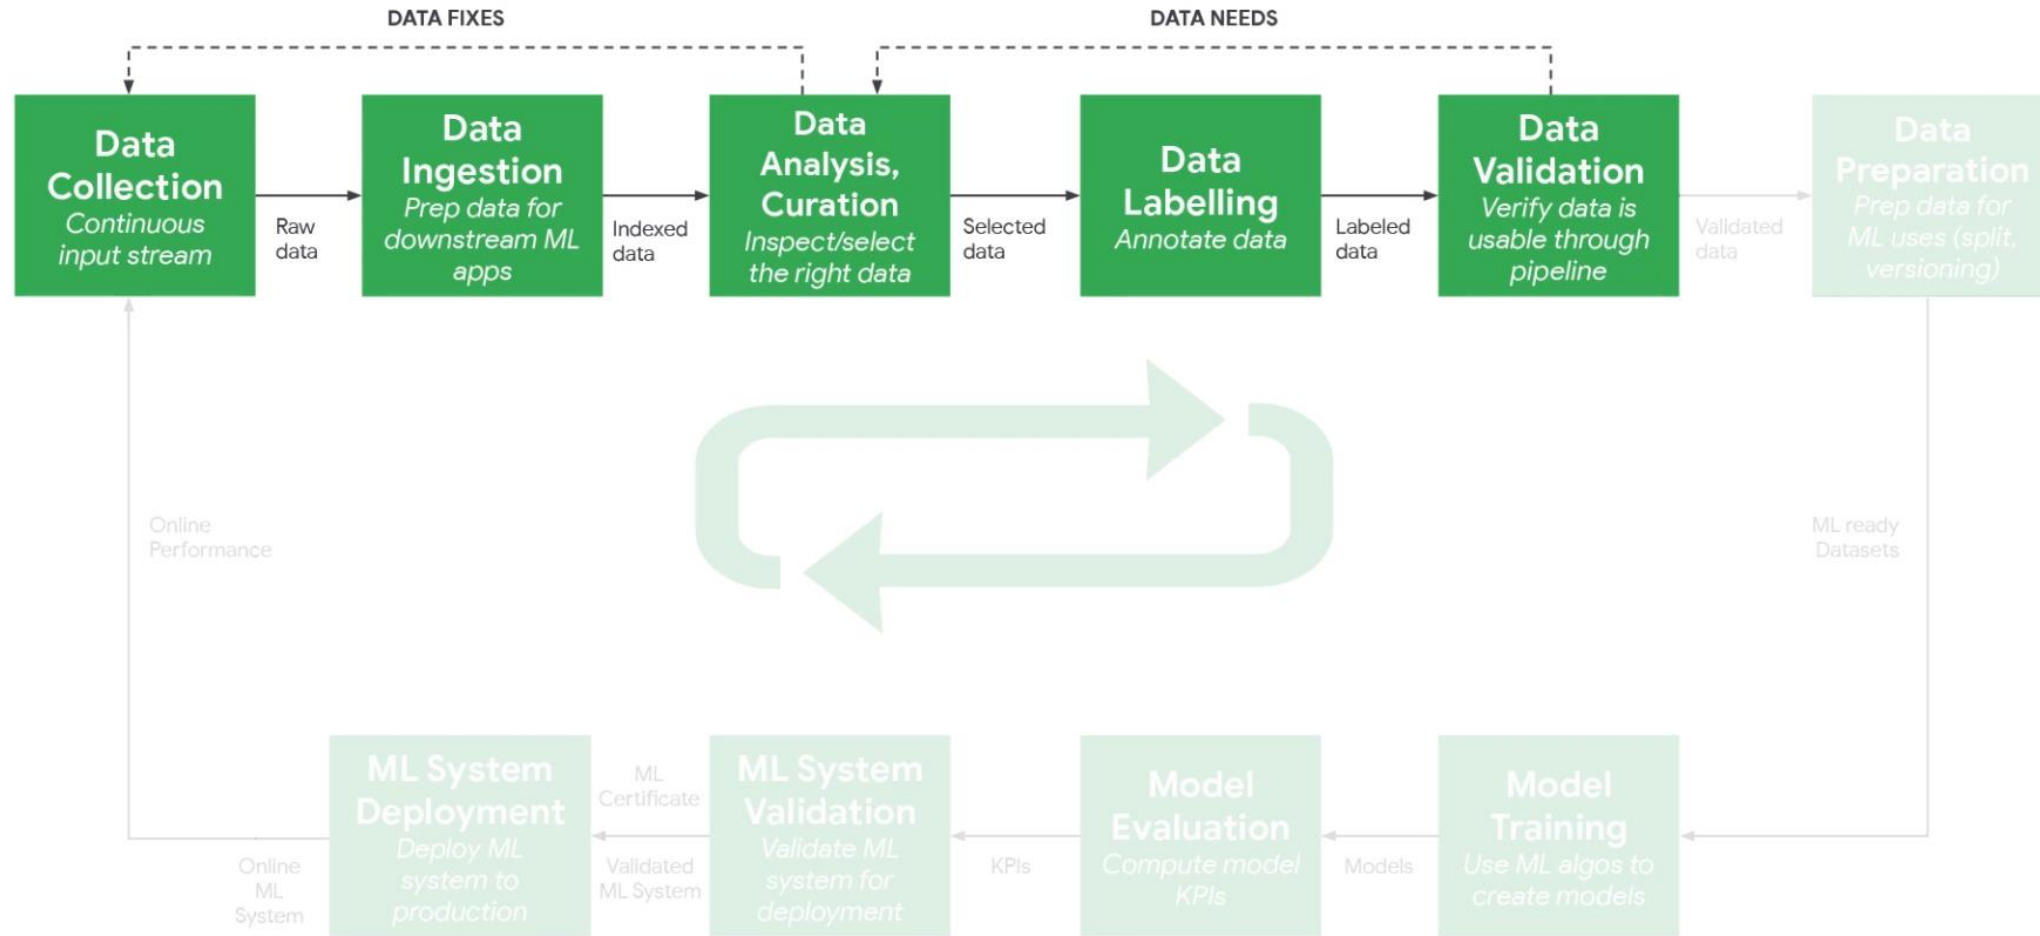
# ML real life cycle



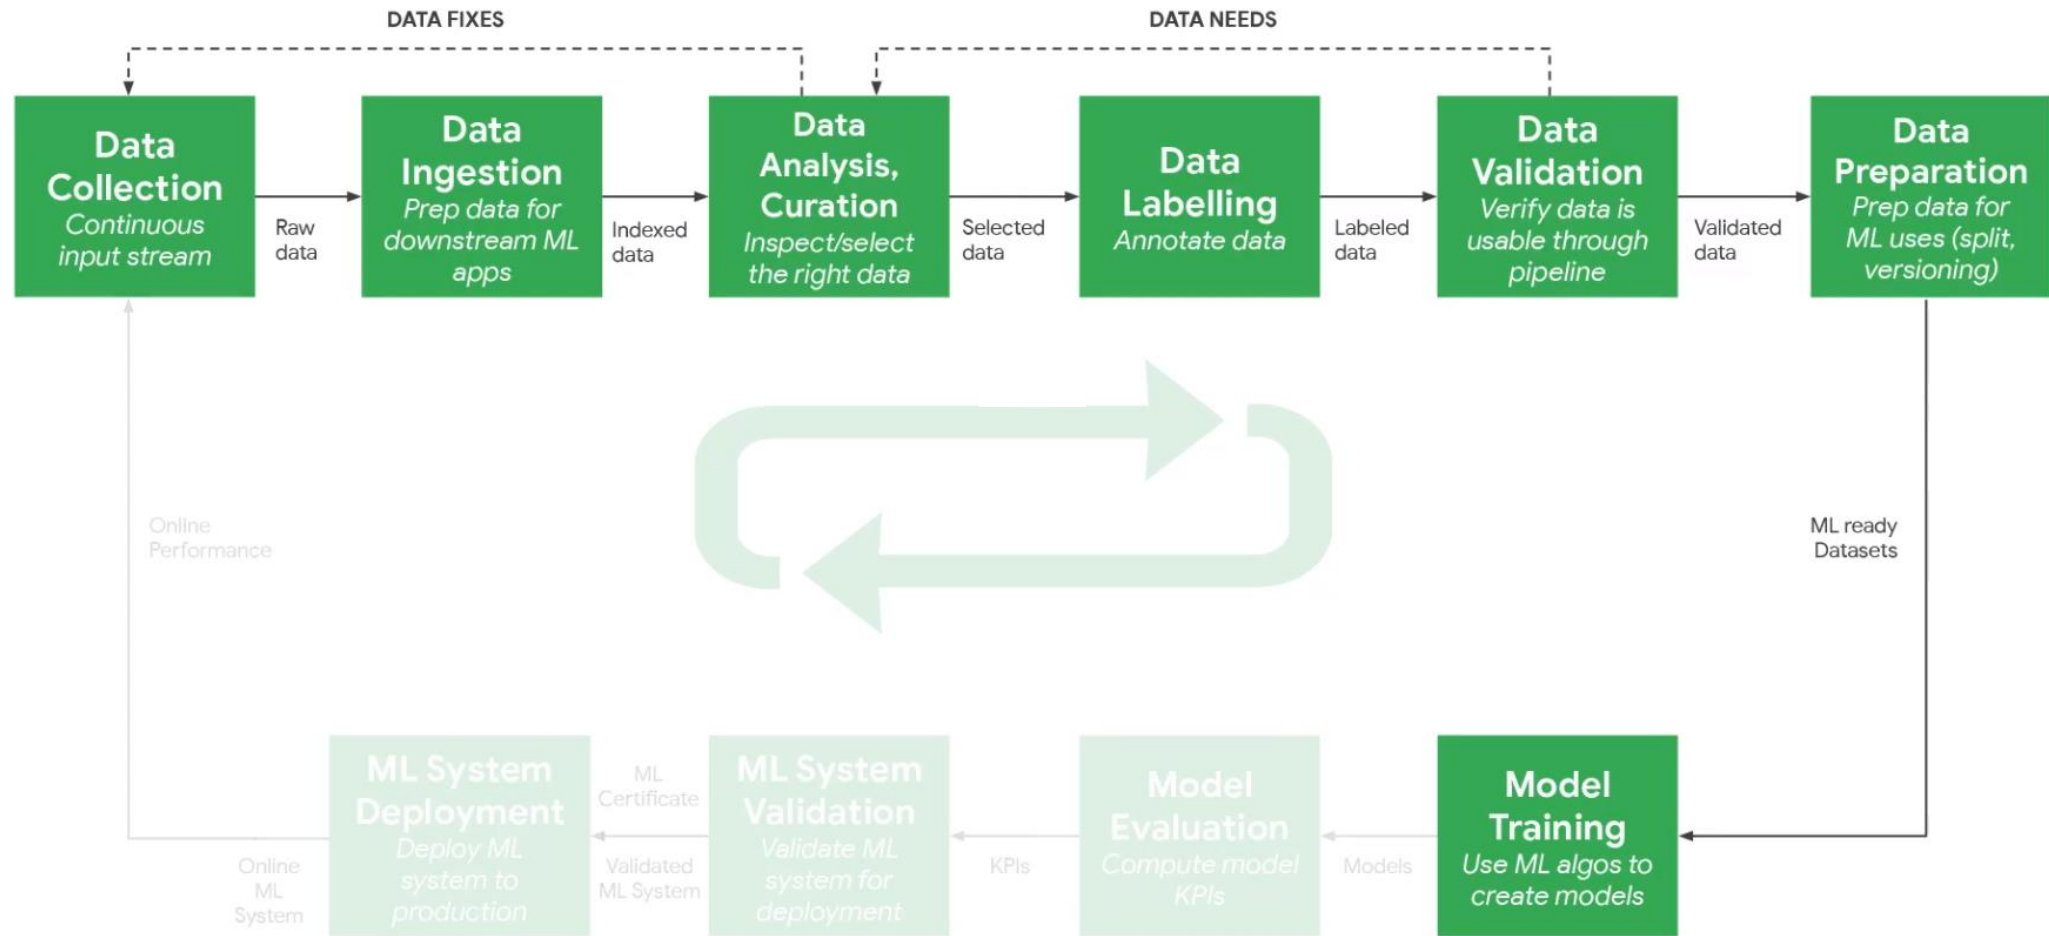
# ML real life cycle



# ML real life cycle

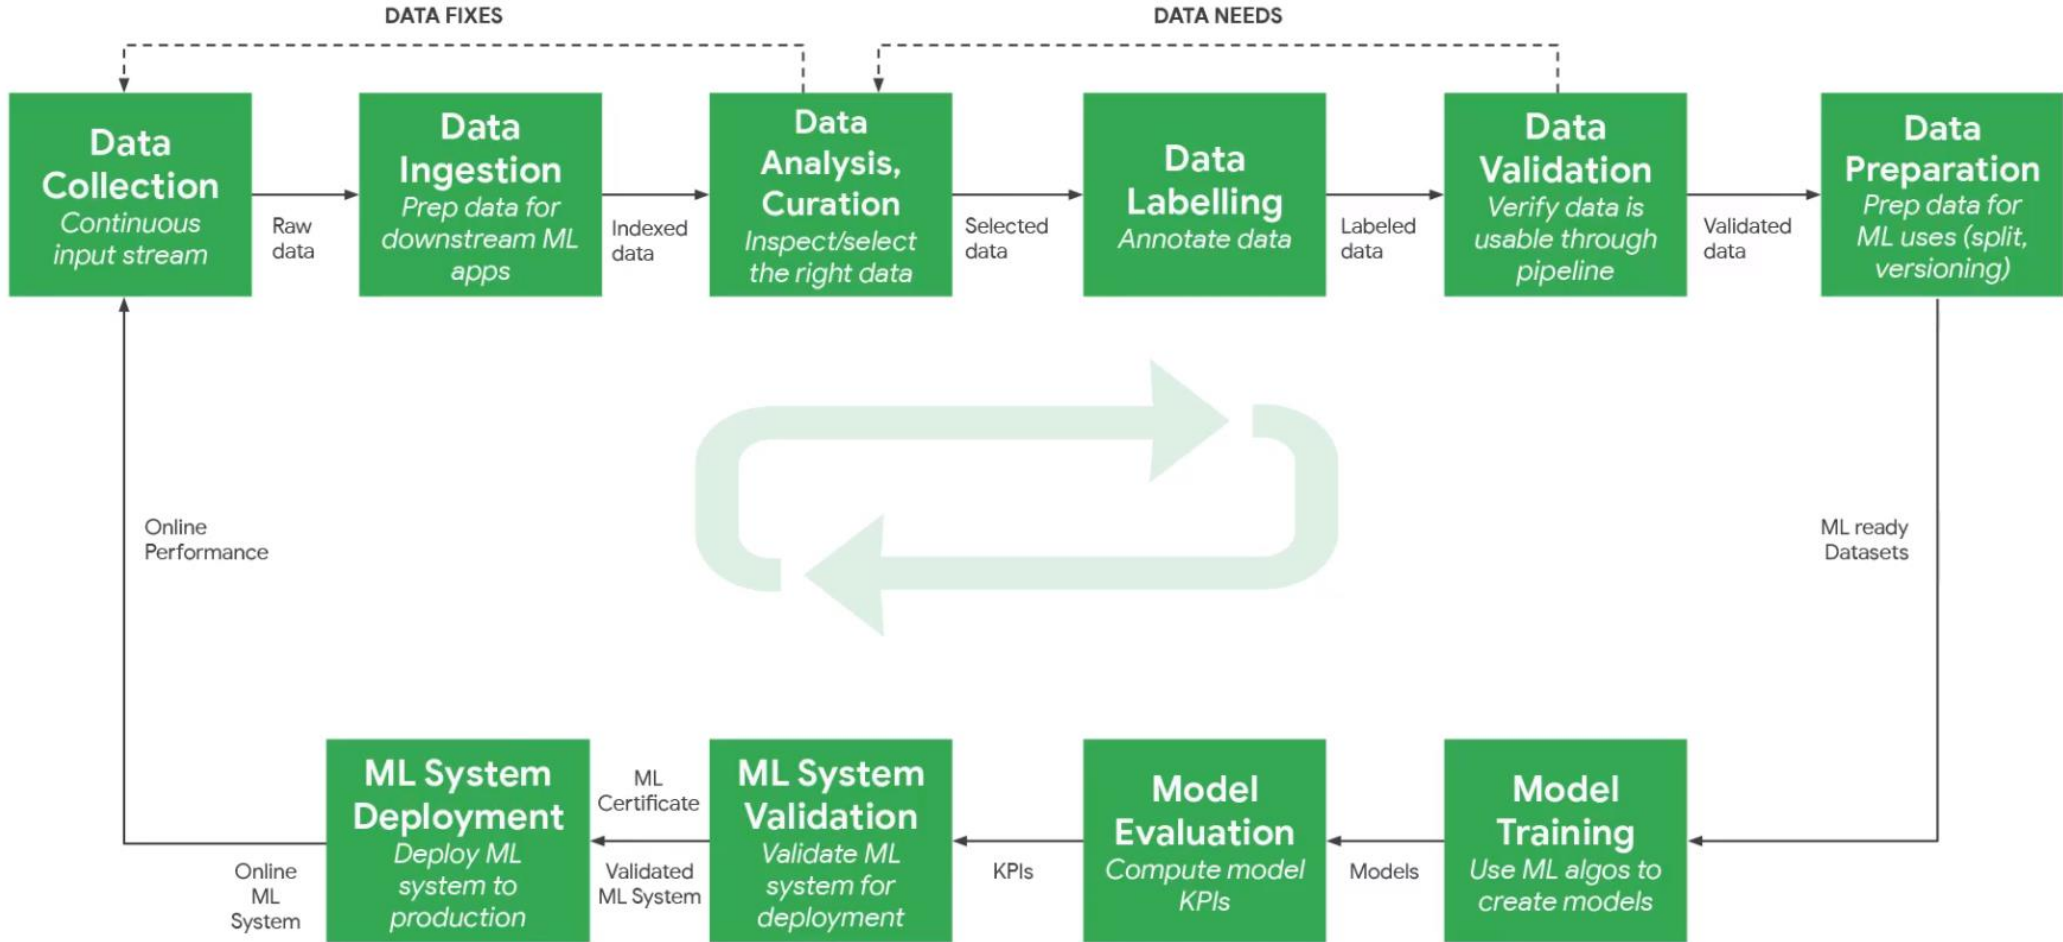


# ML real life cycle

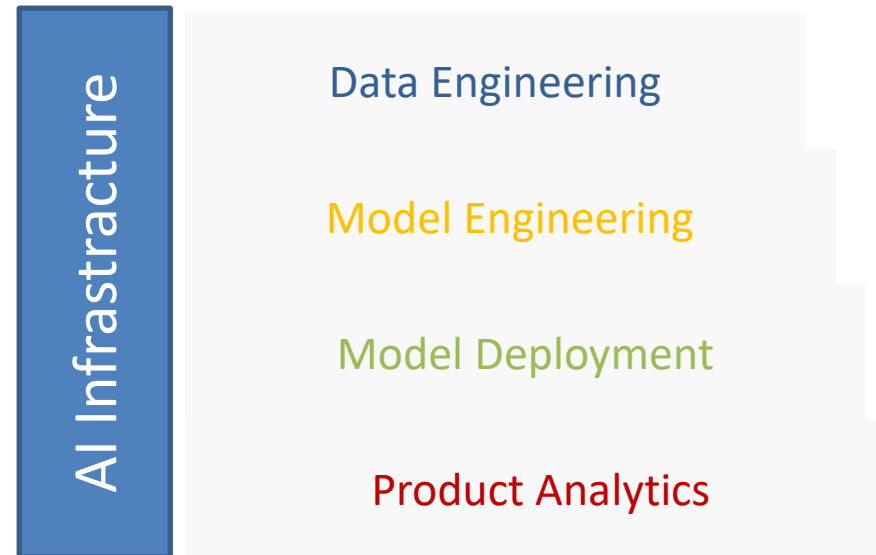




# ML real life cycle



# (Tiny) Machine Learning Workflow



# Data Engineering

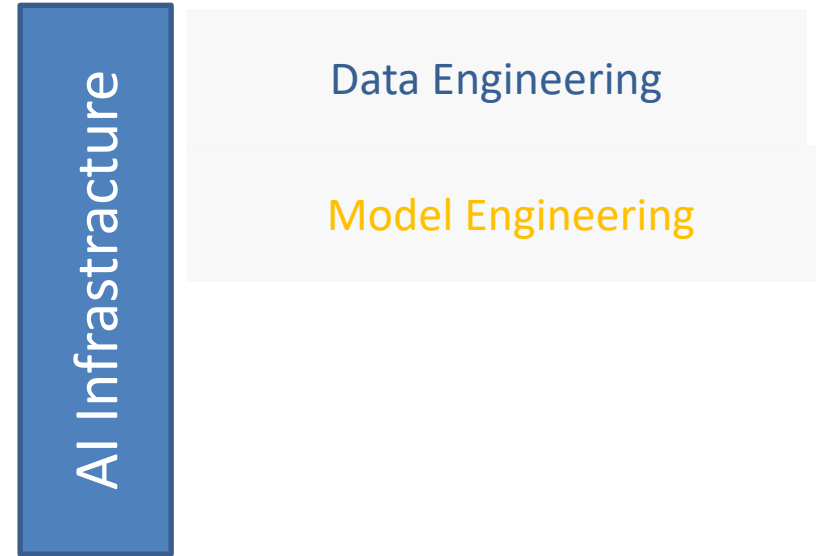
- Defining data **requirements**
- **Collecting** data
- **Labelling** the data
- Inspect and **clean** the data
- Prepare data for **training**
- **Augment** the data
- **Add more data**

AI Infrastructure

Data Engineering

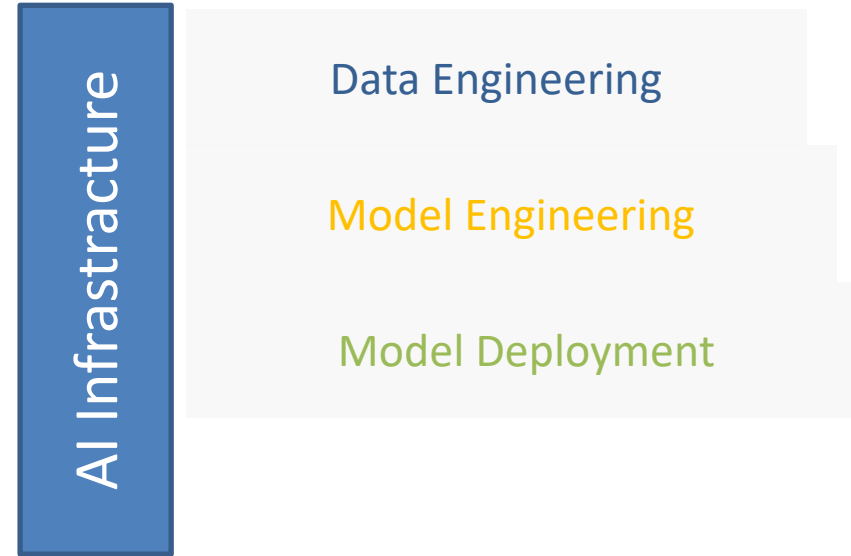
# Model Engineering

- Training ML models
- Improving training **speed**
- Setting target **metrics**
- Evaluating against metrics
- **Optimizing** model training
- Keeping up with **SOTA**



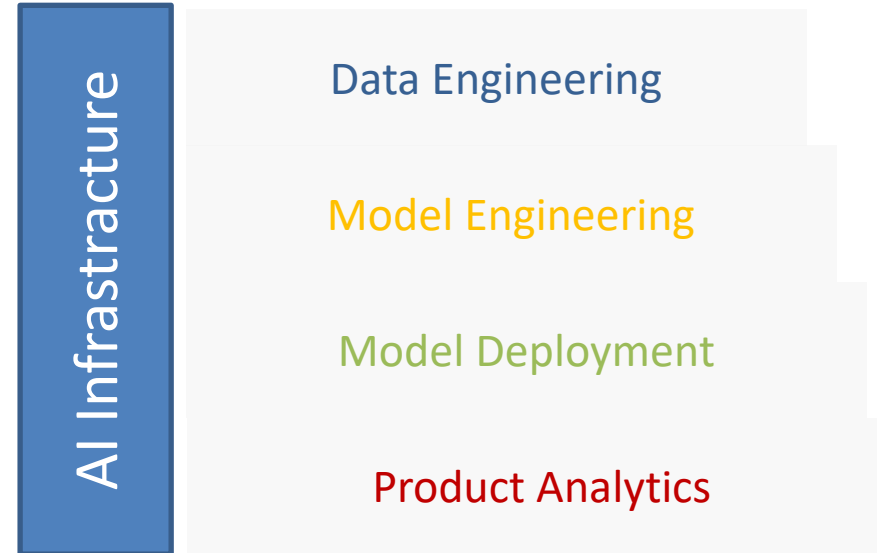
# Model Deployment

- Model conversion
- Performance optimization
- Energy-aware optimizations
- Security and privacy
- Inference serving APIs
- On-device fine-tuning

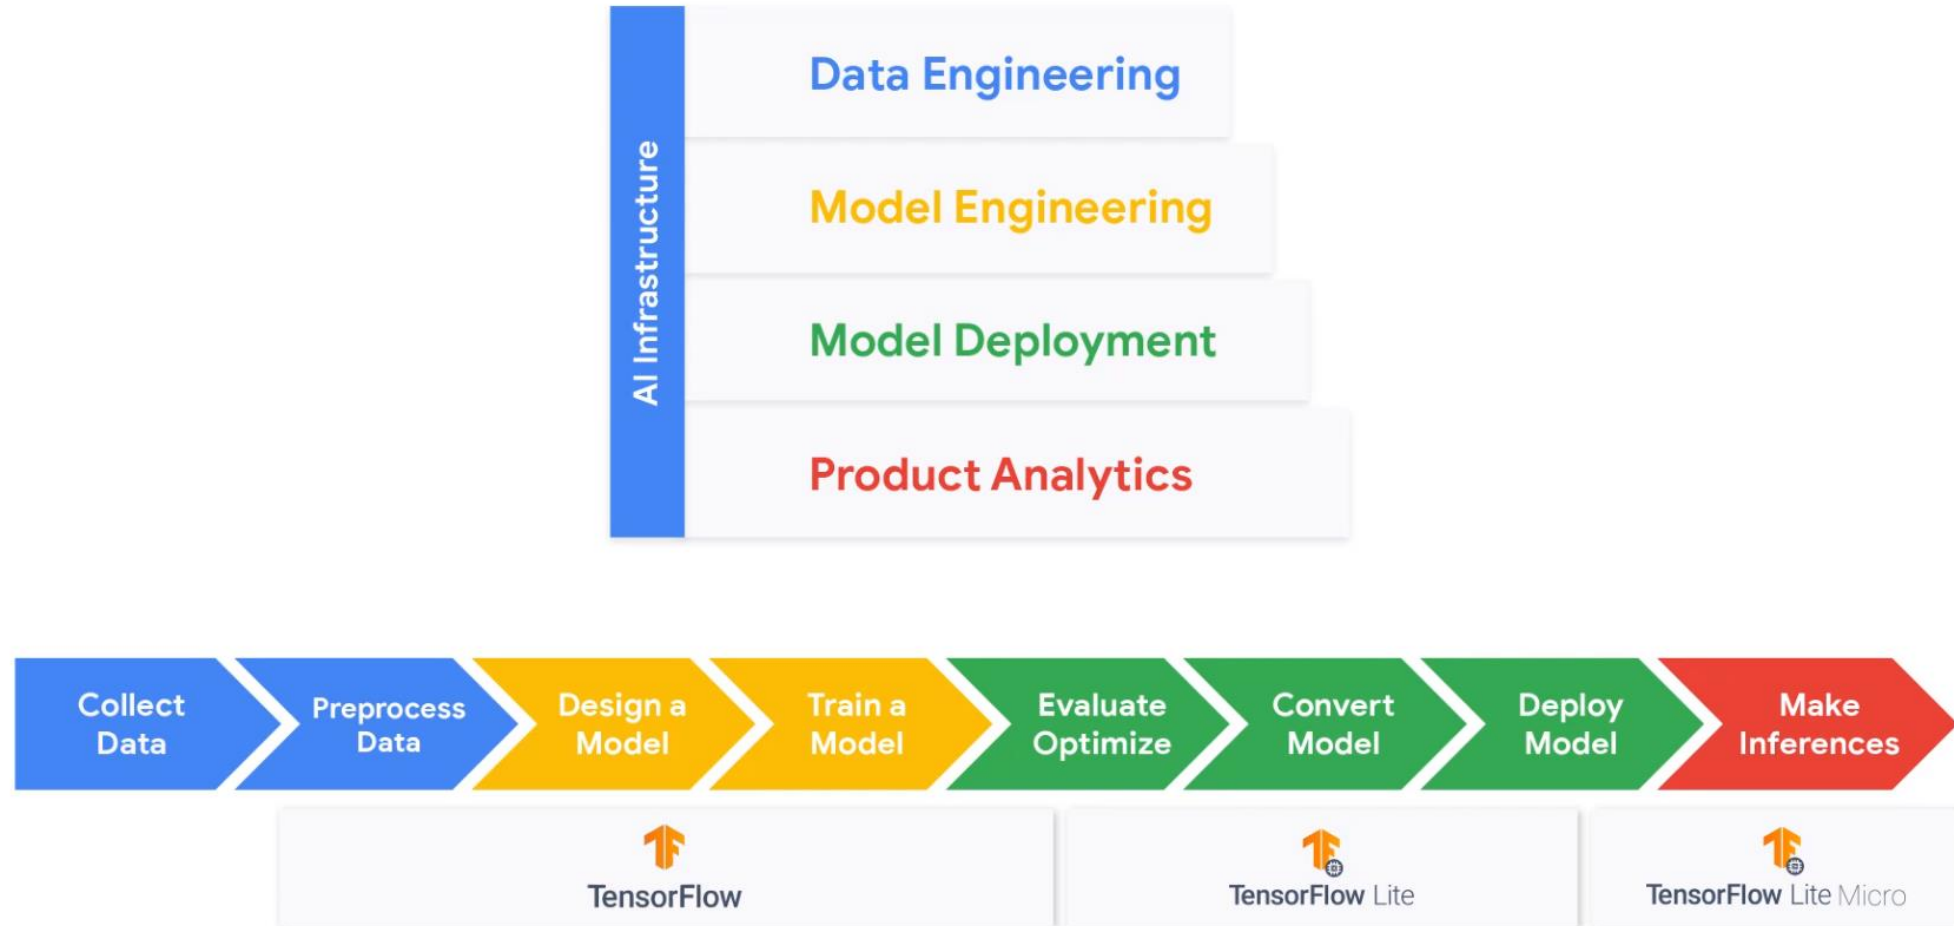


# Product Analysis

- Dashboards
- Field data evaluation
- **Value-added** for business
- Opportunities for **advancement** and improvements



# TinyML workflow






# Differences between targets and operative frameworks

	 TensorFlow	 TensorFlow Lite	 TensorFlow Lite Micro
Training	Yes	No	No
Inference	Yes <i>(but inefficient on edge)</i>	Yes <i>(and efficient)</i>	Yes <i>(and even more efficient)</i>
How Many Ops	~1400	~130	~50
Native Quantization Tooling + Support	No	Yes	Yes



	 TensorFlow	 TensorFlow Lite	 TensorFlow Lite Micro
Needs an OS	Yes	Yes	No
Memory Mapping of Models	No	Yes	Yes
Delegation to accelerators	Yes	Yes	No

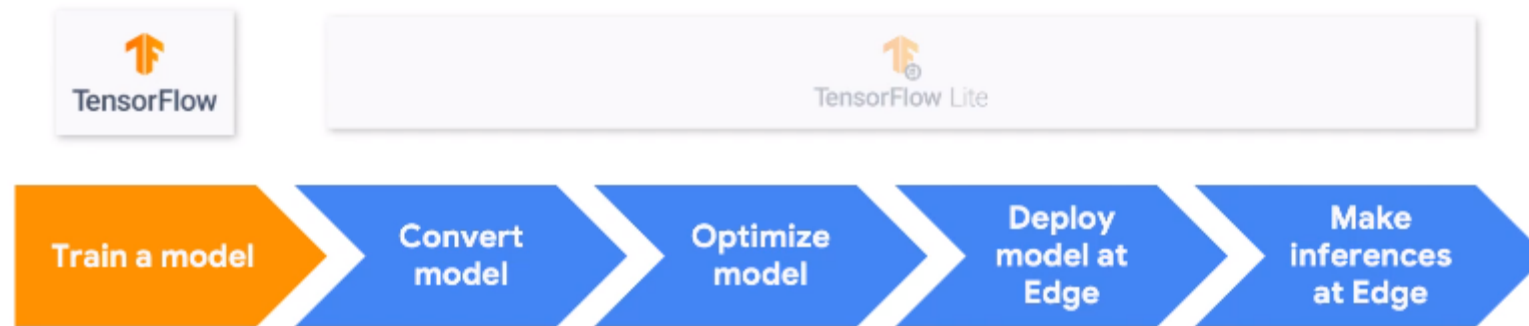
	 TensorFlow	 TensorFlow Lite	 TensorFlow Lite Micro
Base Binary Size	3MB+	100KB	~10 KB
Base Memory Footprint	~5MB	300KB	20KB
Optimized Architectures	X86, TPUs, GPUs	Arm Cortex A, x86	Arm Cortex M, DSPs, MCUs



**POLITECNICO**  
MILANO 1863

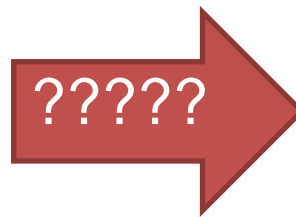
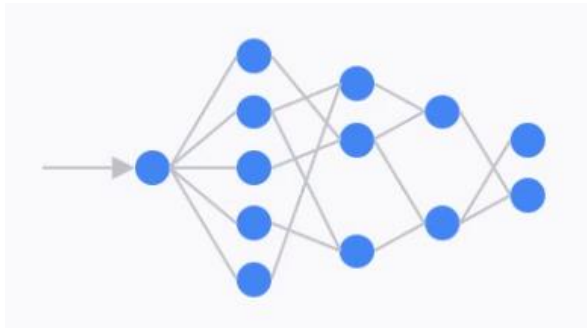
# Edge Devices

# Let's start with EDGE



# A comparison

- MobileNet (2015)
  - MobileNetv1
    - 70.6% top-1 accuracy
    - 16.9MB in size
- Arduino Nano BLE 33 sense
  - Has only 256KB of RAM!
  - 1 MB of flash



# Using the TFLITE converter

- Export saved model:
- Use the TFLite converter:
- Save the TFLite model:
- Create TFLite interpreter:
- Test inference:

```
export_dir = 'saved_model/1'  
tf.saved_model.save(model, export_dir)
```

```
converter = tf.lite.TFLiteConverter.from_saved_model(export_dir)  
tflite_model = converter.convert()
```

```
import pathlib  
tflite_model_file = pathlib.Path('model.tflite')  
tflite_model_file.write_bytes(tflite_model)
```

```
interpreter = tf.lite.Interpreter(model_path=tflite_model_file)
```

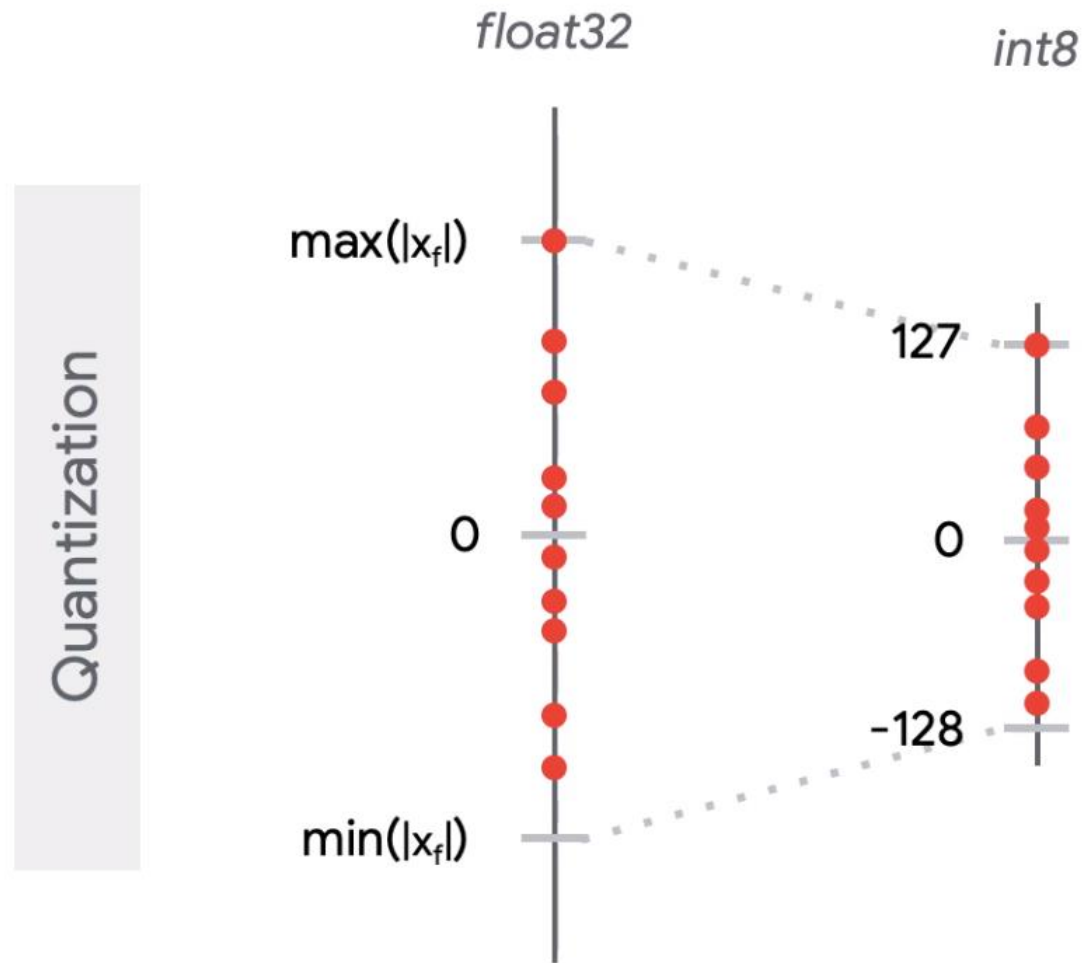
```
# Get input and output tensors.  
input_details = interpreter.get_input_details()  
output_details = interpreter.get_output_details()  
to_predict = # Input data in the same shape as what the model expects  
interpreter.set_tensor(input_details[0]['index'], to_predict)  
  
tflite_results = interpreter.get_tensor(output_details[0]['index'])
```

# Using the TFLite converter

COLAB:

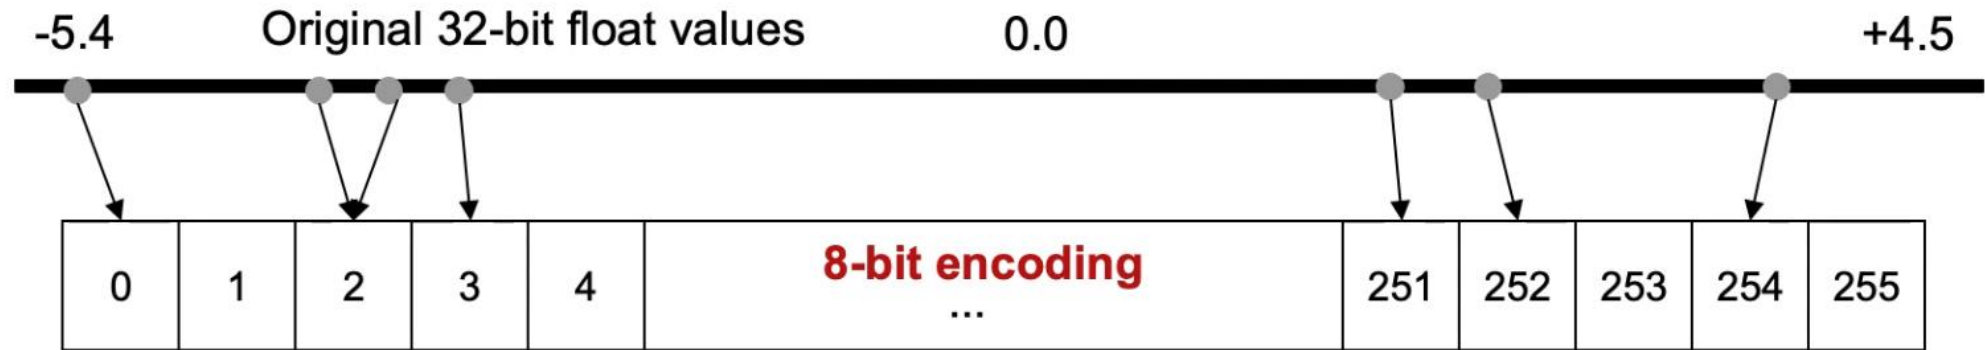
<https://colab.research.google.com/github/tinyMLx/colabs/blob/master/3-3-7-RunningTFLiteModels.ipynb>

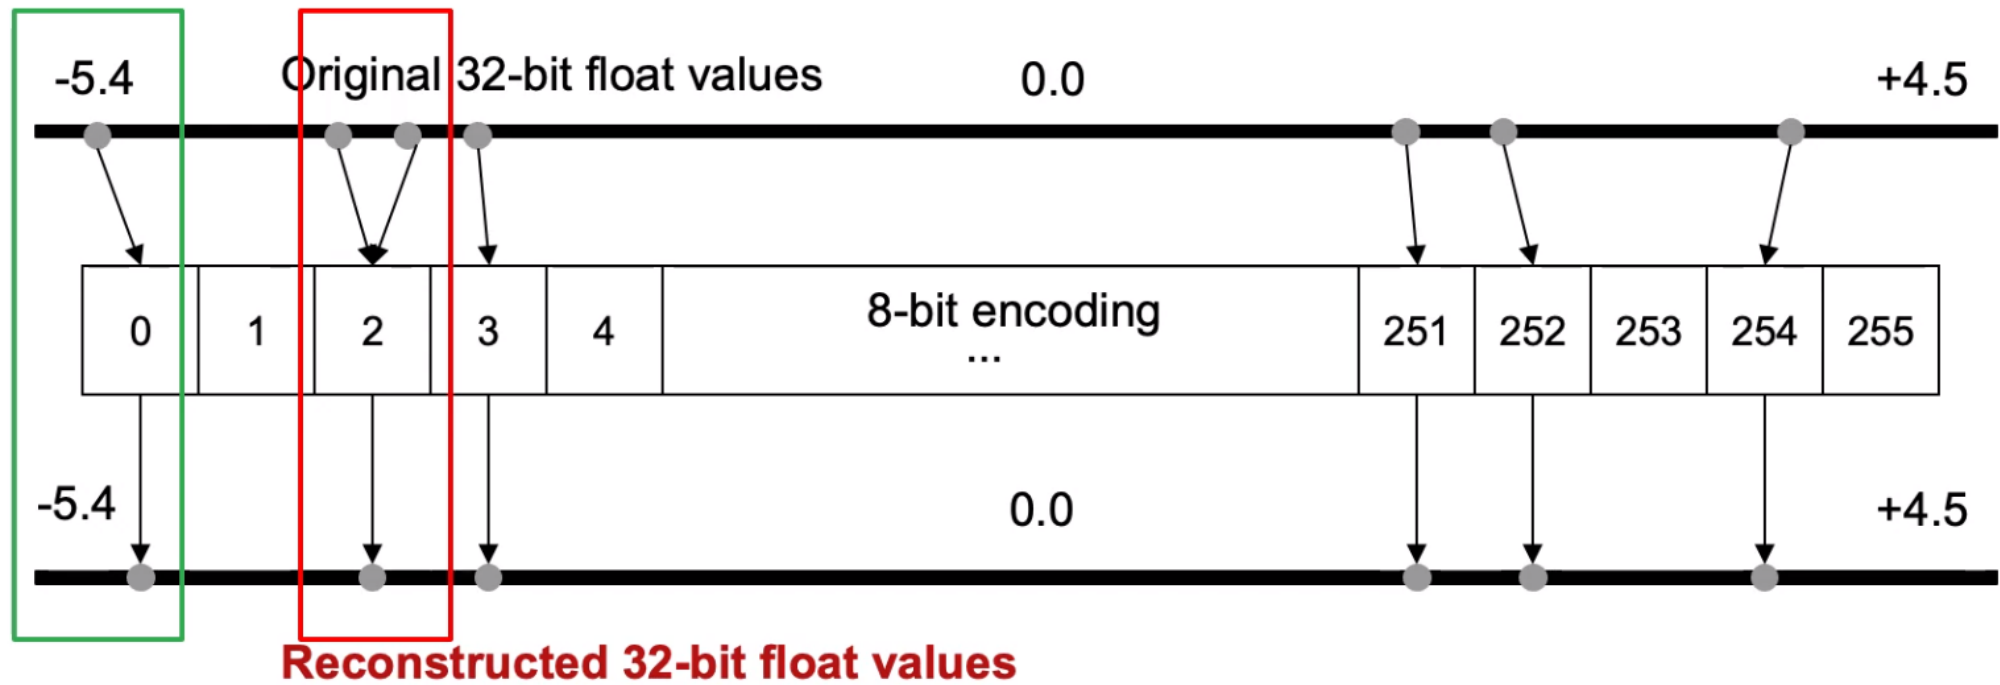
# Quantization





# Quantization

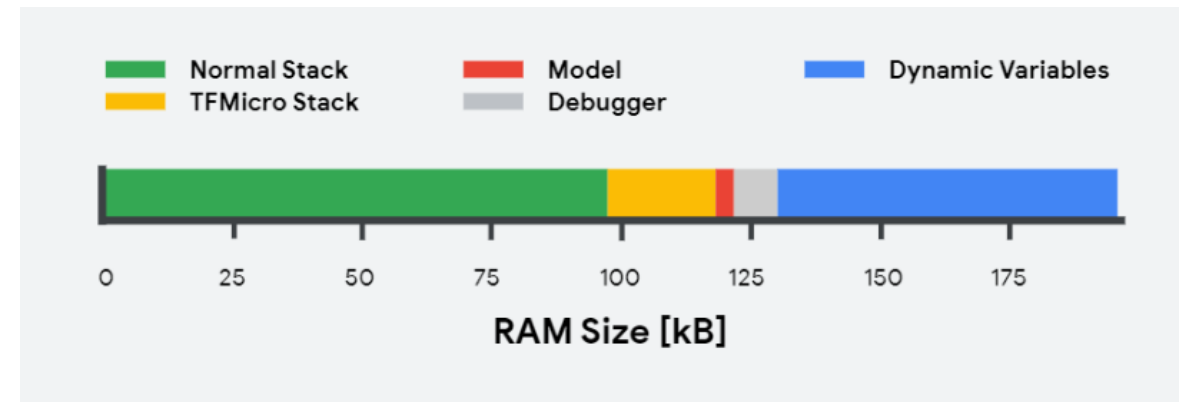
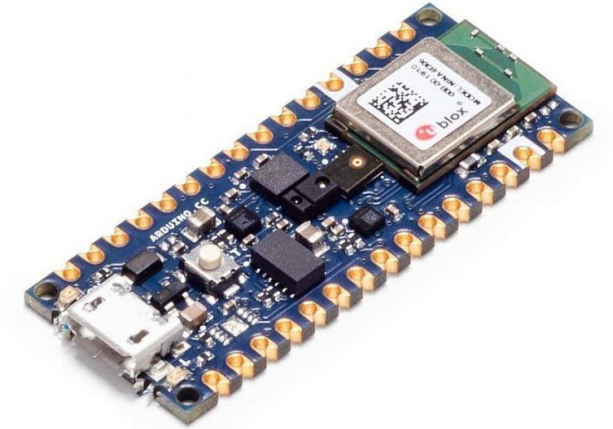




# Why quantizing is (almost) always a good idea

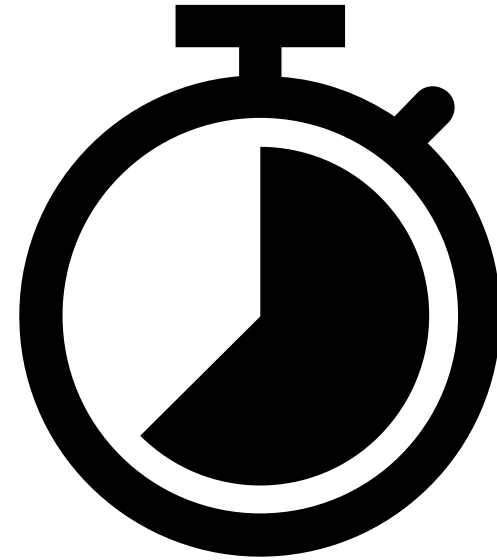
- **Storage and Memory size:**  
when quantizing we have an immediate  $\approx 4x$  reduction on the storage Memory, and a 2/4x reduction in peak RAM usage depending on the type of quantization you choose

- Arduino Nano BLE 33 sense has 1 MB of flash memory



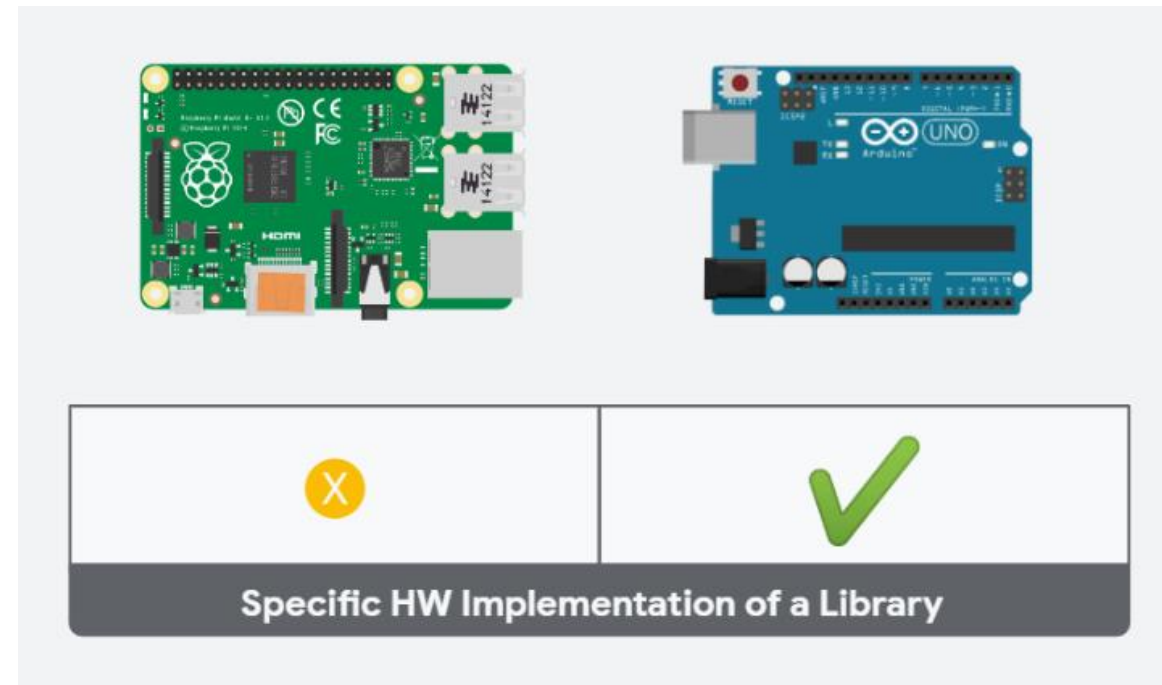
# Why quantizing is (almost) always a good idea

- **Inference Latency**: by performing integer arithmetics instead of floating point arithmetics all the computations on the device are much faster. The expected reduction in time is in the order of 2/4x
- **Energy saving**: for the same reasons performing 8bit operations is much less energy-hungry than 32bit floating point arithmetics.



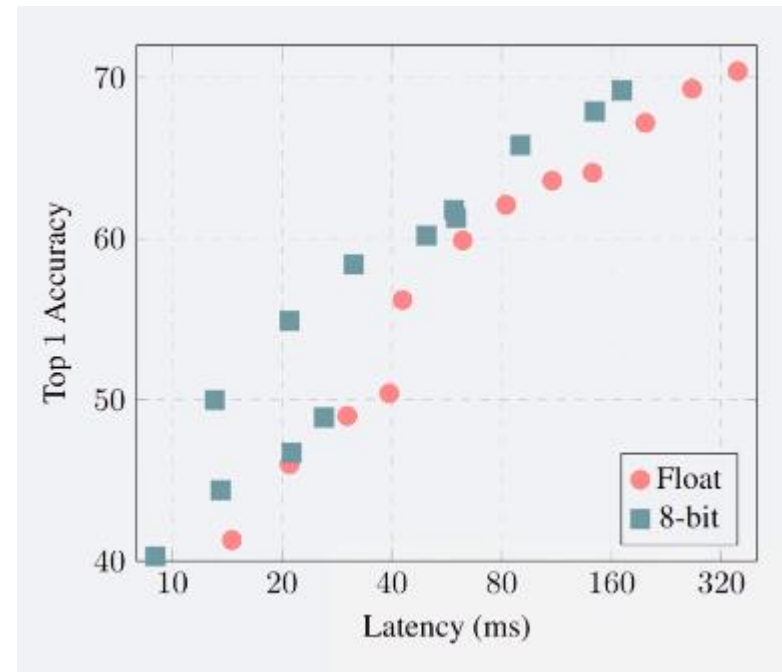
# Why quantizing is (almost) always a good idea

- **Portability:** not every device can perform floating point 32 bits arithmetics. Everyone instead can perform 8bit integer arithmetics.



# Accuracy-latency(memory) trade-off

Quantization works well but performance can suffer of **accuracy-loss** during inference.



# Enabling conversions optimization

Default optimization:

```
converter.optimizations = [tf.lite.Optimize.DEFAULT]

tflite_model = converter.convert()
tflite_model_file = 'converted_model.tflite'

with open(tflite_model_file, "wb") as f:
    f.write(tflite_model)
```

Quantization with  
representative  
dataset optimization:

```
converter.optimizations = [tf.lite.Optimize.DEFAULT]

def representative_data_gen():
    for input_value, _ in test_batches.take(100):
        yield [input_value]

converter.representative_dataset = representative_data_gen
converter.target_spec.supported_ops = [tf.lite.OpsSet.TFLITE_BUILTINS_INT8]

tflite_model = converter.convert()
tflite_model_file = 'converted_model.tflite'

with open(tflite_model_file, "wb") as f:
    f.write(tflite_model)
```

# Full integer optimization

```
def representative_data_gen():
    for input_value in tf.data.Dataset.from_tensor_slices(train_images).batch(1).take(100):
        yield [input_value]

converter = tf.lite.TFLiteConverter.from_keras_model(model)
converter.optimizations = [tf.lite.Optimize.DEFAULT]
converter.representative_dataset = representative_data_gen
# Ensure that if any ops can't be quantized, the converter throws an error
converter.target_spec.supported_ops = [tf.lite.OpsSet.TFLITE_BUILTINS_INT8]
# Set the input and output tensors to uint8 (APIs added in r2.3)
converter.inference_input_type = tf.uint8
converter.inference_output_type = tf.uint8

tflite_model_quant = converter.convert()
```

In case input and output type == uint8, remember to quantize the input before feeding it to the network! (and de-quantize the output, in case you need it)

```
# Check if the input type is quantized, then rescale input data to uint8
if input_details['dtype'] == np.uint8:
    input_scale, input_zero_point = input_details["quantization"]
    test_image = test_image / input_scale + input_zero_point
```



# Optimizing the network

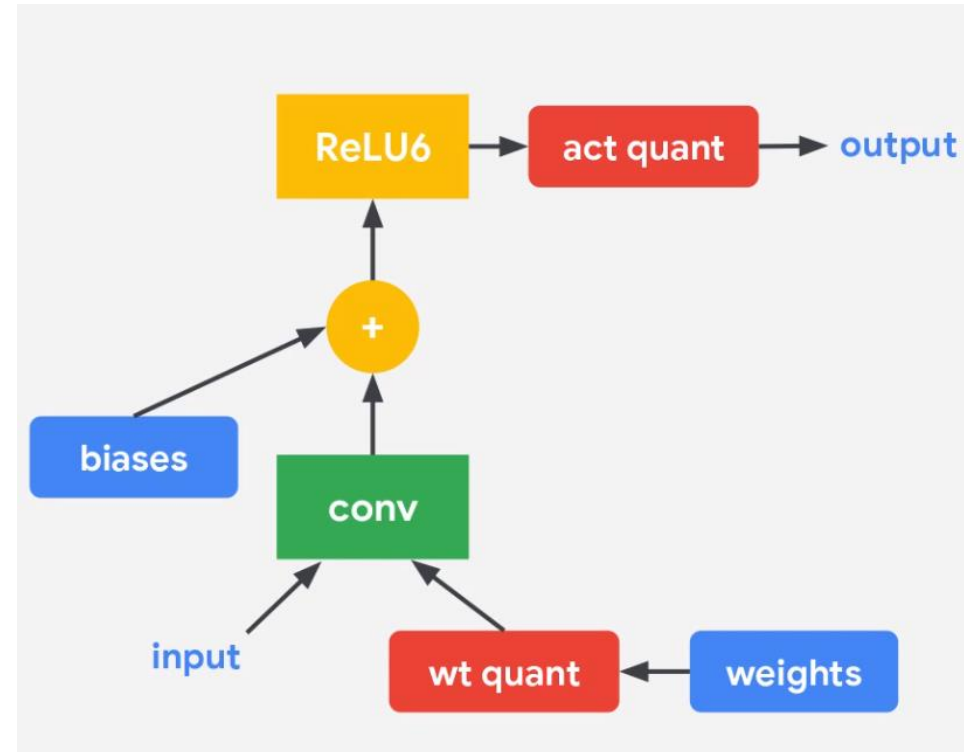
COLAB:

<https://colab.research.google.com/github/tinyMLx/colabs/blob/master/3-3-10-TFLiteOptimizations.ipynb#scrollTo=0RTZmndkcZFP>

# Quantization-aware training

Quantization aware training emulates inference-time quantization during training, creating a model that downstream tools will use to produce actually quantized models. The quantized models use lower-precision (e.g. 8-bit instead of 32-bit float), leading to benefits during deployment.

- Mimic the inference path during the training phase.
- Expose the training pipeline to the errors introduced by quantization
- Allow the training phase to recover the error «naturally»



# Results comparison

	Floating-point Baseline	Post-training Quantization (PTQ)	Quantization-Aware Training (QAT)
MobileNet v1 1.0 224	71.03%	69.57%	71.06%
MobileNet v2 1.0 224	70.77%	70.20%	70.01%
Resnet v1 50	76.30%	75.95%	76.10%

# How to quantization aware-training

```
import tensorflow_model_optimization as tfmot

quantize_model = tfmot.quantization.keras.quantize_model

# q_aware stands for for quantization aware.
q_aware_model = quantize_model(model)

# `quantize_model` requires a recompile.
q_aware_model.compile(optimizer='adam',
                      loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
                      metrics=['accuracy'])
```

# Quantization aware training

COLAB:

<https://colab.research.google.com/github/tinyMLx/colabs/blob/master/3-3-12-QAT.ipynb#scrollTo=w7fztWsAOHTz>

# Hands on: rock paper scissor

<https://colab.research.google.com/drive/1vAXuU9bDbD90W6fnpx1crvbQrzpnOcaU?usp=sharing>



**POLITECNICO**  
MILANO 1863

# Appendix

# Credits and reference

- “TinyML: Machine Learning with TensorFlow Lite on Arduino and Ultra-Low-Power Microcontrollers”, Daniel Situnayake, Pete Warden, O'Reilly Media, Inc.
- Online course:
  - <https://www.edx.org/professional-certificate/harvardx-tiny-machine-learning>
- A lot more material on TinyML:
  - <http://tinymml.seas.harvard.edu/>